A Content-Driven Micro-Video Recommendation Dataset at Scale

CONTENTS

01

Motivation

Motivation (MoRec)

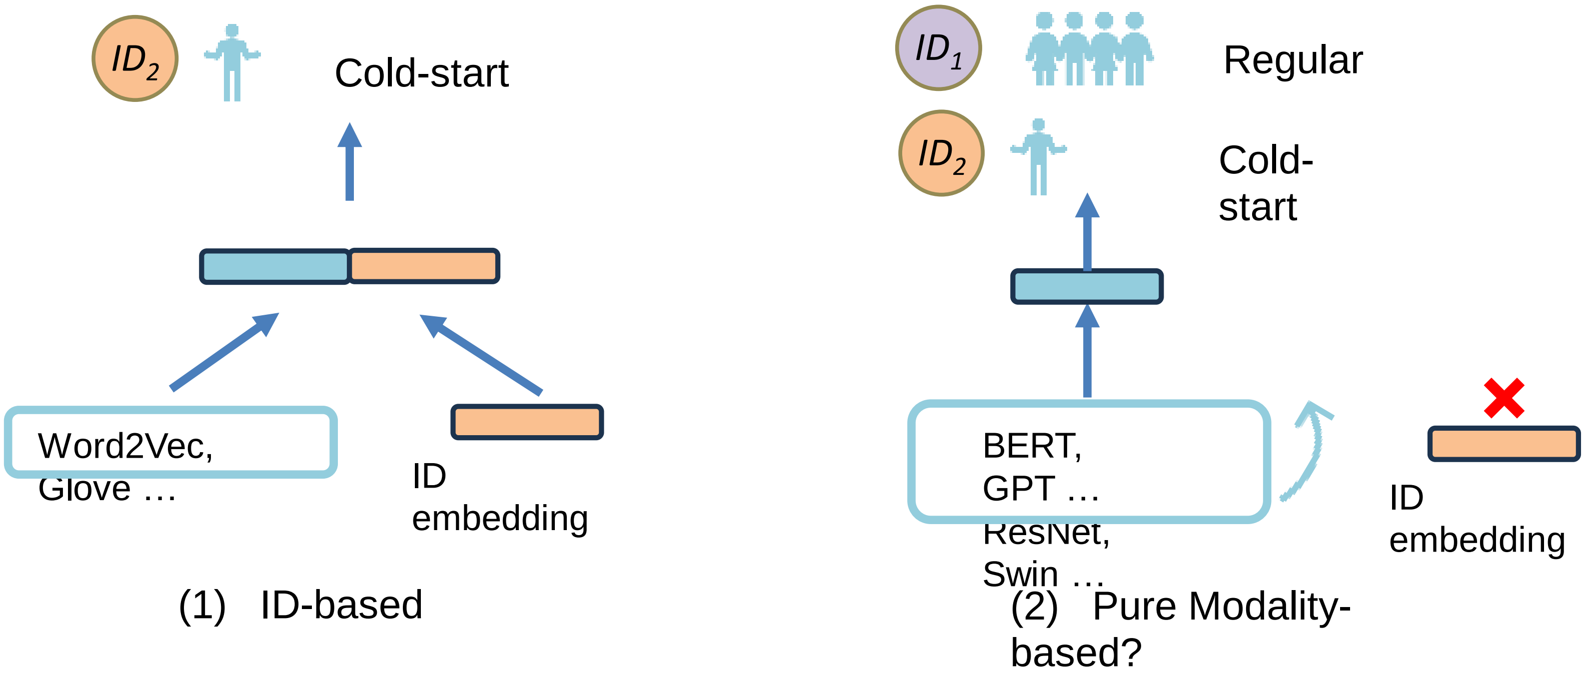

Motivation (MoRec)

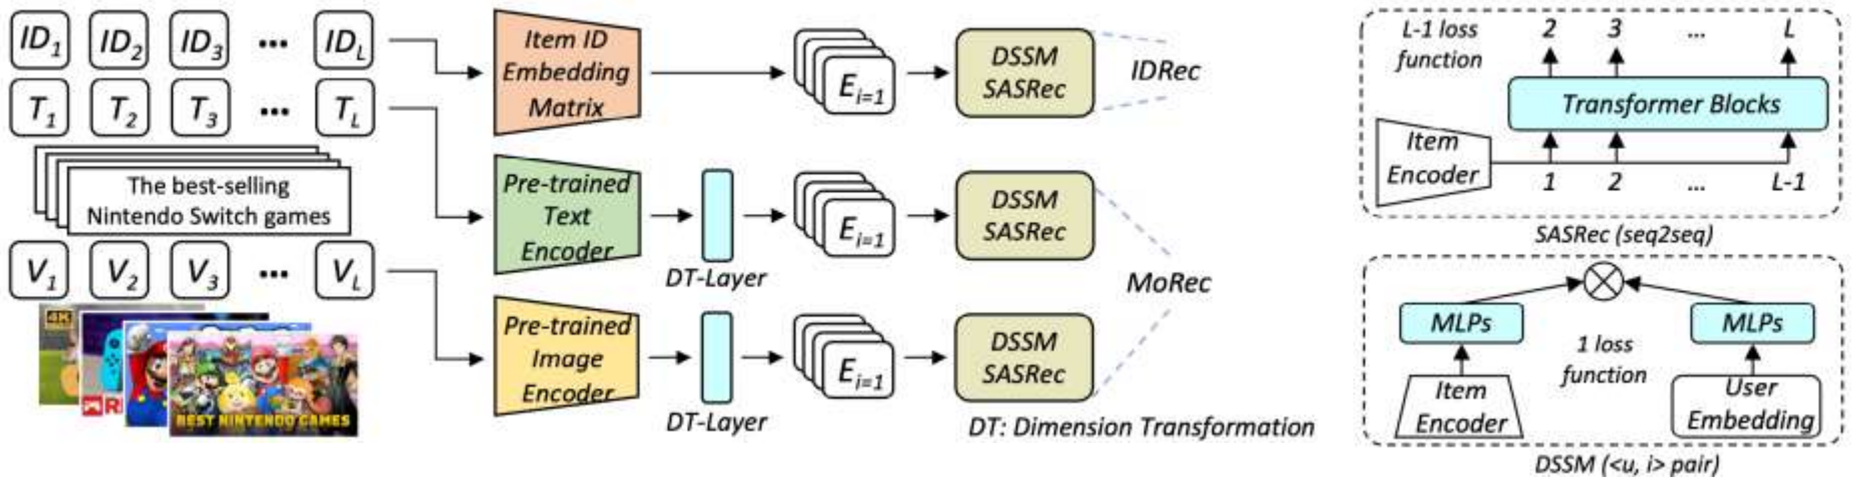

Nowadays, with the help of current textual/visual encoders, MoRec can be comparable to or even better than IDRec

Motivation (MoRec)

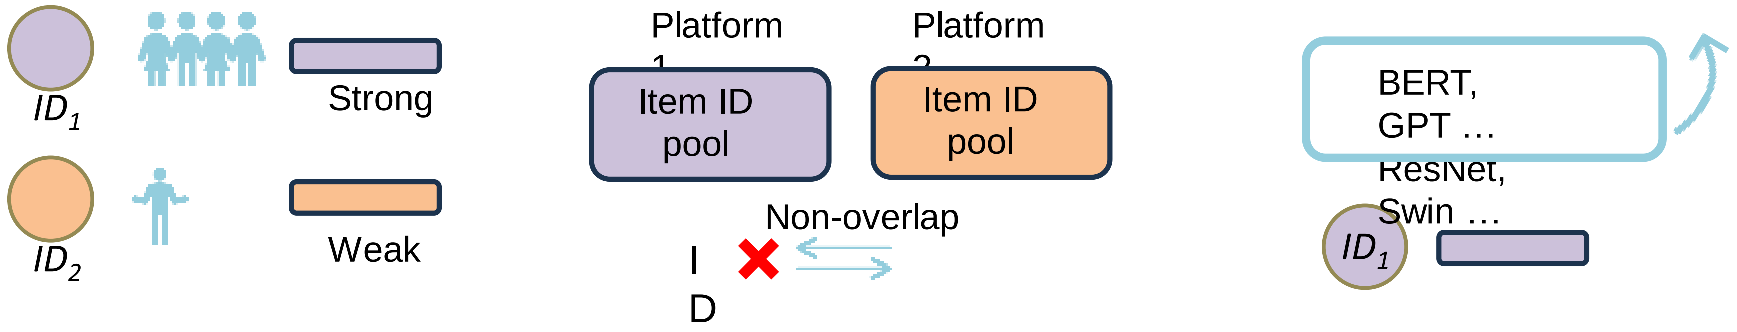

(1) Cold-start setting

(2) Transfer

(3) Benefit from CV/NLP/MM

Motivation (Future Direction)

- Modality-based Recommendation
- Micro-video Recommendation
- Foundation Models for Recommender Systems
- “one4all” Paradigm

Motivation (Lack of Datasets)

- Domain
- Raw Content
- Scale
- Modality Diversity

02

MicroLens

MicroLens (Dataset)

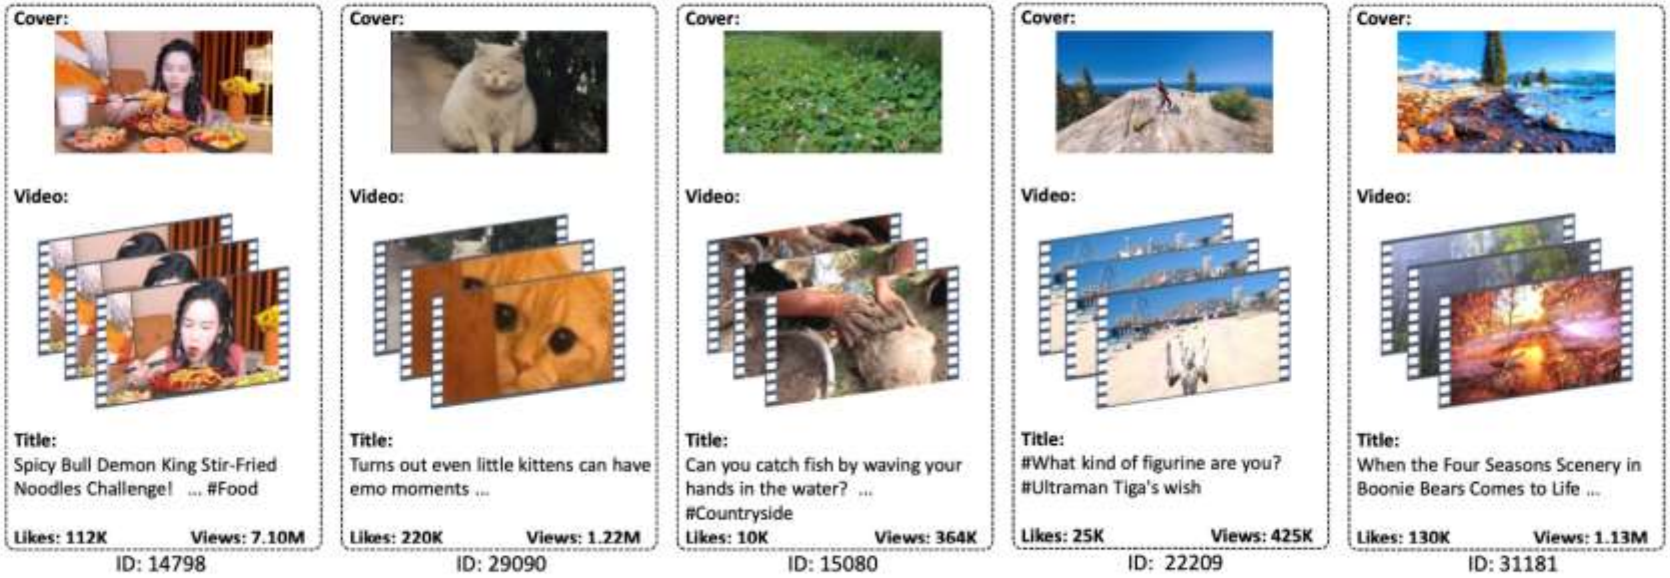

Figure 2: Item examples in MicroLens.

MicroLens (Experiments)

- VideoRec
  - End-to-end manner
  - Train recommender model and video encoder simultaneously
- Investigate how *RS* benefits from *Video Understanding*
- 3 recommender models
  - CNN-based (NextItNet)
  - RNN-based (GRU4Rec)
  - Transformer-based (SASRec)
- 15 video encoders
  - R3D-r18, X3D-xs, C2D-r50, I3D- r50, X3D-s, Slow-r50, X3D-m, R3D-r50, SlowFast-r50, CSN-r101, X3D-l, SlowFast-r101, MViT-B-16x4, MViT-B-32x3, and VideoMAE

03

Findings

Findings (Benchmark Results)

Class	Model	HR@10	NDCG@10	HR@20	NDCG@20
IDRec (CF)	DSSM [29]	0.0394	0.0193	0.0654	0.0258
	LightGCN [26]	0.0372	0.0177	0.0618	0.0239
	NFM [25]	0.0313	0.0159	0.0480	0.0201
	DeepFM [17]	0.0350	0.0170	0.0571	0.0225
IDRec (SR)	NexItNet [62]	0.0805	0.0442	0.1175	0.0535
	GRU4Rec [27]	0.0782	0.0423	0.1147	0.0515
	SASRec [31]	0.0909	0.0517	0.1278	0.0610
VIDRec (Frozen Encoder)	$\text{YouTube}_{\text{ID}}$	0.0461	0.0229	0.0747	0.0301
	$\text{YouTube}_{\text{ID+V}}$ [7]	0.0392	0.0188	0.0648	0.0252
	MMGCN_{ID}	0.0141	0.0065	0.0247	0.0092
	$\text{MMGCN}_{\text{ID+V}}$ [54]	0.0214	0.0103	0.0374	0.0143
	GRCN_{ID}	0.0282	0.0131	0.0497	0.0185
	$\text{GRCN}_{\text{ID+V}}$ [53]	0.0306	0.0144	0.0547	0.0204
	$\text{DSSM}_{\text{ID+V}}$	0.0279	0.0137	0.0461	0.0183
	$\text{SASRec}_{\text{ID+V}}$	0.0799	0.0415	0.1217	0.0520
VideoRec (E2E Learning)	$\text{NexItNet}_{\text{V}}$ [62]	0.0862	0.0466	0.1246	0.0562
	$\text{GRU4Rec}_{\text{V}}$ [27]	0.0954	0.0517	0.1377	0.0623
	SASRec_{V} [31]	0.0948	0.0515	0.1364	0.0619

Findings (Benchmark Results)

- Methods: IDRec, VIDRec and VideoRec
  - We do not search parameters exhaustively for VideoRec
  - Only 5 frames of each video were used
- Findings： raw video content > pre-extracted frozen features

Findings (Video Understanding Meets Recommender Systems)

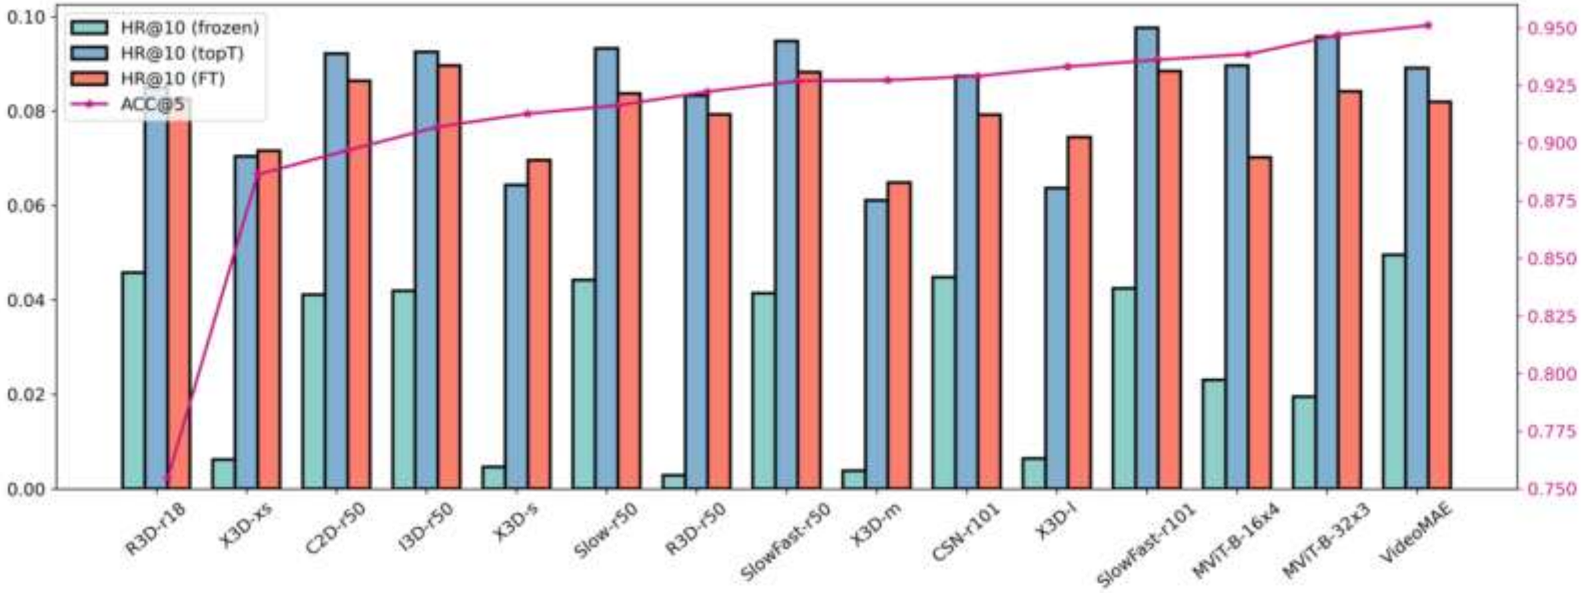

Figure 4: Video recommendation accuracy (bar charts) vs. video classification accuracy (purple line). Frozen means that the video encoder is fixed without parameter update, topT means that only the top few layers of the video encoder are fine-tuned, and FT means full parameters are fine-tuned.

Findings (Video Understanding Meets Recommender Systems)

Table 6: Performance of VideoRec with 15 video encoders. "Pretrain Settings" are the adopted frame length and sample rate from the pre-trained checkpoint. ACC@5 is the accuracy in the video classification task.

Model	Architecture	Depth	Pretrain Settings	ACC@5	HR@10 (frozen)	NDCG@10 (frozen)	HR@10 (topT)	NDCG@10 (topT)	HR@10 (FT)	NDCG@10 (FT)
R3D-r18 [47]	ResNet	R18	16x4	75.45	4.58	2.56	8.50	4.48	7.50	3.48
X3D-xs [10]	Xception	XS	4x12	88.63	0.62	0.33	7.04	3.57	6.04	2.57
C2D-r50 [52]	ResNet	R50	8x8	89.68	4.11	2.27	9.22	4.88	8.22	3.88
I3D-r50 [4]	ResNet	R50	8x8	90.70	4.19	2.36	9.25	5.01	8.25	4.01
X3D-s [10]	Xception	S	13x6	91.27	0.47	0.24	6.43	3.25	5.43	2.25
Slow-r50 [8]	ResNet	R50	8x8	91.63	4.42	2.42	9.32	4.99	8.33	3.99
X3D-m [10]	Xception	M	16x5	92.72	0.38	0.20	6.11	3.13	5.11	2.13
R3D-r50 [47]	ResNet	R50	16x4	92.23	0.28	0.14	8.33	4.34	7.33	3.34
SlowFast-r50 [11]	ResNet	R50	8x8	92.69	4.14	2.35	9.48	5.15	8.48	4.15
CSN-r101 [46]	ResNet	R101	32x2	92.90	4.48	2.52	8.74	4.71	7.74	3.71
X3D-l [10]	Xception	L	16x5	93.31	0.64	0.34	6.37	3.32	5.37	2.32
SlowFast-r101 [11]	ResNet	R101	16x8	93.61	4.25	2.36	**9.76**	**5.3**	**8.76**	**4.31**
MViT-B-16x4 [9]	VIT	B	16x4	93.85	2.30	1.33	8.96	4.79	7.96	3.79
MViT-B-32x3 [9]	VIT	B	32x3	94.69	1.95	1.11	9.57	5.11	8.57	4.11
VideoMAE [45]	Transformer	VIT-B	16x4	**95.10**	**4.96**	**2.76**	8.91	4.77	7.91	3.77

Findings (Video Understanding Meets Recommender Systems)

- Better CV performance ≠ Higher recommendation accuracy
    - E.g., the worst video classification model R3D-r18
- In RS, finetuning top layers > full finetuning
    - full finetuning the video encoders is not necessary in recommender systems

Findings (Video Understanding Meets Recommender Systems)

- Knowledge learned from video understanding helps video recommendation
- Video semantic representations learned from CV task are not universal
  - a linear layer is not enough produce the same results as finetuning

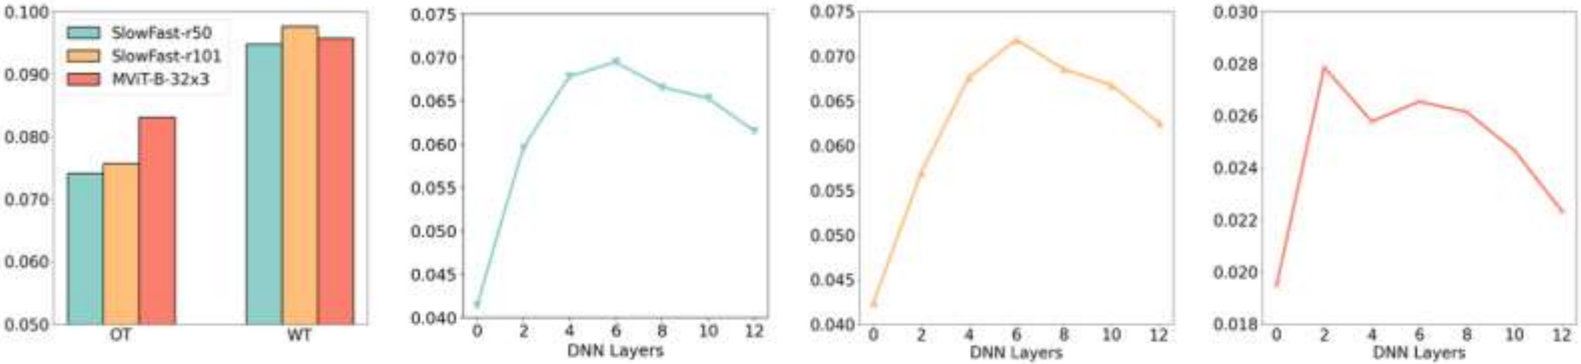

(a) OT v.s. WT (b) SlowFast-r50 (c) SlowFast-r101 (d) MVIT-B-32x3

Figure 5: Ablation study of video encoders. (d) "WT" refers to the video encoders in SASRec$_V$ have pre-trained weights from the video classification task, while "OT" denotes that they are randomly initialized. (b) (c) (d) are performance change by adding DNN layers on top of three frozen encoders.

Findings (Video Understanding Meets Recommender Systems)

- Our study is the first to show that raw video features can potentially replace ID features in both warm and cold item recommendation settings

Table 8: Comparison of VideoRec and IDRec in regular and warm settings using SASRec as the backbone. "Warm-20" denotes that items with less than 20 interactions were removed from the original MicroLens-100K.

Model	Regular		Warm-20		Warm-50		Warm-200	
	H@10	N@10	H@10	N@10	H@10	N@10	H@10	N@10
IDRec	0.0909	0.0517	0.1068	0.0615	0.6546	0.4103	0.7537	0.4412
SlowFast-r101	0.0976	0.0531	0.1130	0.0606	0.7458	0.4463	0.8482	0.4743
MViT-B-32x3	0.0957	0.0511	0.1178	0.0639	0.7464	0.4530	0.9194	0.4901
SlowFast-r50	0.0948	0.0515	0.1169	0.0642	0.7580	0.4614	0.8141	0.4870

Findings (Video Understanding Meets Recommender Systems)

- Our study is the first to show that raw video features can potentially replace ID features in both warm and cold item recommendation settings

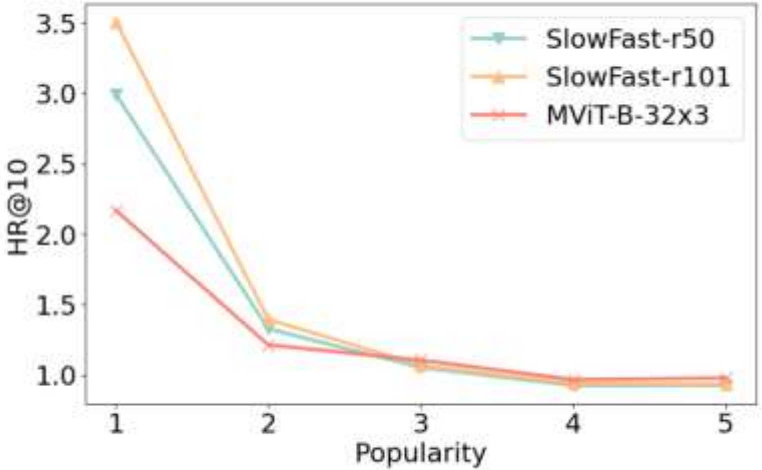

Figure 6: Results in different cold-start scenarios, with the y-axis representing the relative improvement of HR@10, calculated as the ratio of VideoRec to IDRec. The x-axis represents item groups divided by popularity level, the larger number indicates that items in the group are more popular.

Findings (Video Understanding Meets Recommender Systems)

- Summary: This work has taken a key step towards the goal of a universal "one-for-all" recommender paradigm
  - Dataset Support
  - VideoRec Paradigm Exploration

THANKS

FAQ (Collecting)

Why models that adopts MoRec/VideoRec/MMRec paradigms are transferable across platforms?

Why topT is better that FT?

Is there any way to accelerate the training?

Any other usage of MicroLens?